

model_resisc45

April 12, 2023

1 Model RESISC 45

```
[ ]: import os
import pandas as pd
import cv2
import matplotlib.pyplot as plt
import numpy as np
import logging
import contextlib
import random
import platform
import torch
from torch.utils.data import Dataset, DataLoader
from torchvision.transforms import ToTensor
import matplotlib.pyplot as plt

logger = logging.getLogger('train')
logger.setLevel(logging.INFO)

print(platform.platform()) # print current platform
```

macOS-13.3.1-arm64-arm-64bit

1.1 Set constants

```
[ ]: DIRPATH = "/Users/cristianion/Desktop/satimg_data/NWPU-RESISC45"

LABELS = [
    'forest',
    'railway_station',
    'tennis_court',
    'basketball_court',
    'river',
    'storage_tank',
    'harbor',
    'terrace',
    'thermal_power_station',
    'golf_course',
```

```

'runway',
'roundabout',
'bridge',
'industrial_area',
'baseball_diamond',
'mobile_home_park',
'overpass',
'church',
'chaparral',
'railway',
'stadium',
'medium_residential',
'sea_ice',
'intersection',
'lake',
'palace',
'airplane',
'cloud',
'sparse_residential',
'airport',
'snowberg',
'parking_lot',
'commercial_area',
'rectangular_farmland',
'island',
'beach',
'circular_farmland',
'dense_residential',
'ship',
'mountain',
'desert',
'freeway',
'meadow',
'wetland',
'ground_track_field',
]

K_FOLDS = 5

```

```
[ ]: print(len(LABELS)) # number of expected labels
```

45

2 Prepare dataset

```
[ ]: filelist = [(f"{DIRPATH}/{f}", LABELS[LABELS.index(f)]) for f in os.  
↳listdir(DIRPATH) if f in LABELS]  
df = pd.DataFrame(filelist, columns=['dirpath', 'label'])
```

```
[ ]: df.head()
```

```
[ ]:
```

	dirpath	label
0	/Users/cristianion/Desktop/sating_data/NWPU-RE...	forest
1	/Users/cristianion/Desktop/sating_data/NWPU-RE...	railway_station
2	/Users/cristianion/Desktop/sating_data/NWPU-RE...	tennis_court
3	/Users/cristianion/Desktop/sating_data/NWPU-RE...	basketball_court
4	/Users/cristianion/Desktop/sating_data/NWPU-RE...	river

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 45 entries, 0 to 44
```

```
Data columns (total 2 columns):
```

#	Column	Non-Null Count	Dtype
0	dirpath	45 non-null	object
1	label	45 non-null	object

```
dtypes: object(2)
```

```
memory usage: 852.0+ bytes
```

Conclusions - 45 directories found corresponding to 45 labels

```
[ ]: data = []  
for i, DIRPATH in enumerate(df["dirpath"]):  
    images = os.listdir(DIRPATH)  
    images = [f"{DIRPATH}/{img}" for img in images]  
    rows = [(img, df['label'][i]) for img in images]  
    data.extend(rows)  
data = pd.DataFrame(data, columns=["imgpath", "label"])
```

```
[ ]: data.head()
```

```
[ ]:
```

	imgpath	label
0	/Users/cristianion/Desktop/sating_data/NWPU-RE...	forest
1	/Users/cristianion/Desktop/sating_data/NWPU-RE...	forest
2	/Users/cristianion/Desktop/sating_data/NWPU-RE...	forest
3	/Users/cristianion/Desktop/sating_data/NWPU-RE...	forest
4	/Users/cristianion/Desktop/sating_data/NWPU-RE...	forest

```
[ ]: data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31500 entries, 0 to 31499
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   imgpath     31500 non-null  object
1   label       31500 non-null  object
dtypes: object(2)
memory usage: 492.3+ KB

```

Conclusions: - 31500 annotated images

```
[ ]: data['label'].value_counts()
```

```

[ ]: label
forest                700
intersection          700
palace                700
airplane              700
cloud                 700
sparse_residential    700
airport               700
snowberg              700
parking_lot           700
commercial_area       700
rectangular_farmland  700
island                700
beach                 700
circular_farmland     700
dense_residential     700
ship                  700
mountain              700
desert                700
freeway               700
meadow                700
wetland               700
lake                  700
sea_ice               700
railway_station       700
medium_residential    700
tennis_court          700
basketball_court      700
river                 700
storage_tank          700
harbor                700
terrace               700
thermal_power_station  700
golf_course           700

```

```

runway          700
roundabout      700
bridge          700
industrial_area 700
baseball_diamond 700
mobile_home_park 700
overpass        700
church          700
chaparral       700
railway         700
stadium         700
ground_track_field 700
Name: count, dtype: int64

```

Conclusions: - Each category of the 45 categories has 700 samples

2.1 Sample subset of images from dataset


- show a subset of the images
- select some random labels (max 5 labels)
- show 5 images for each label to show variance

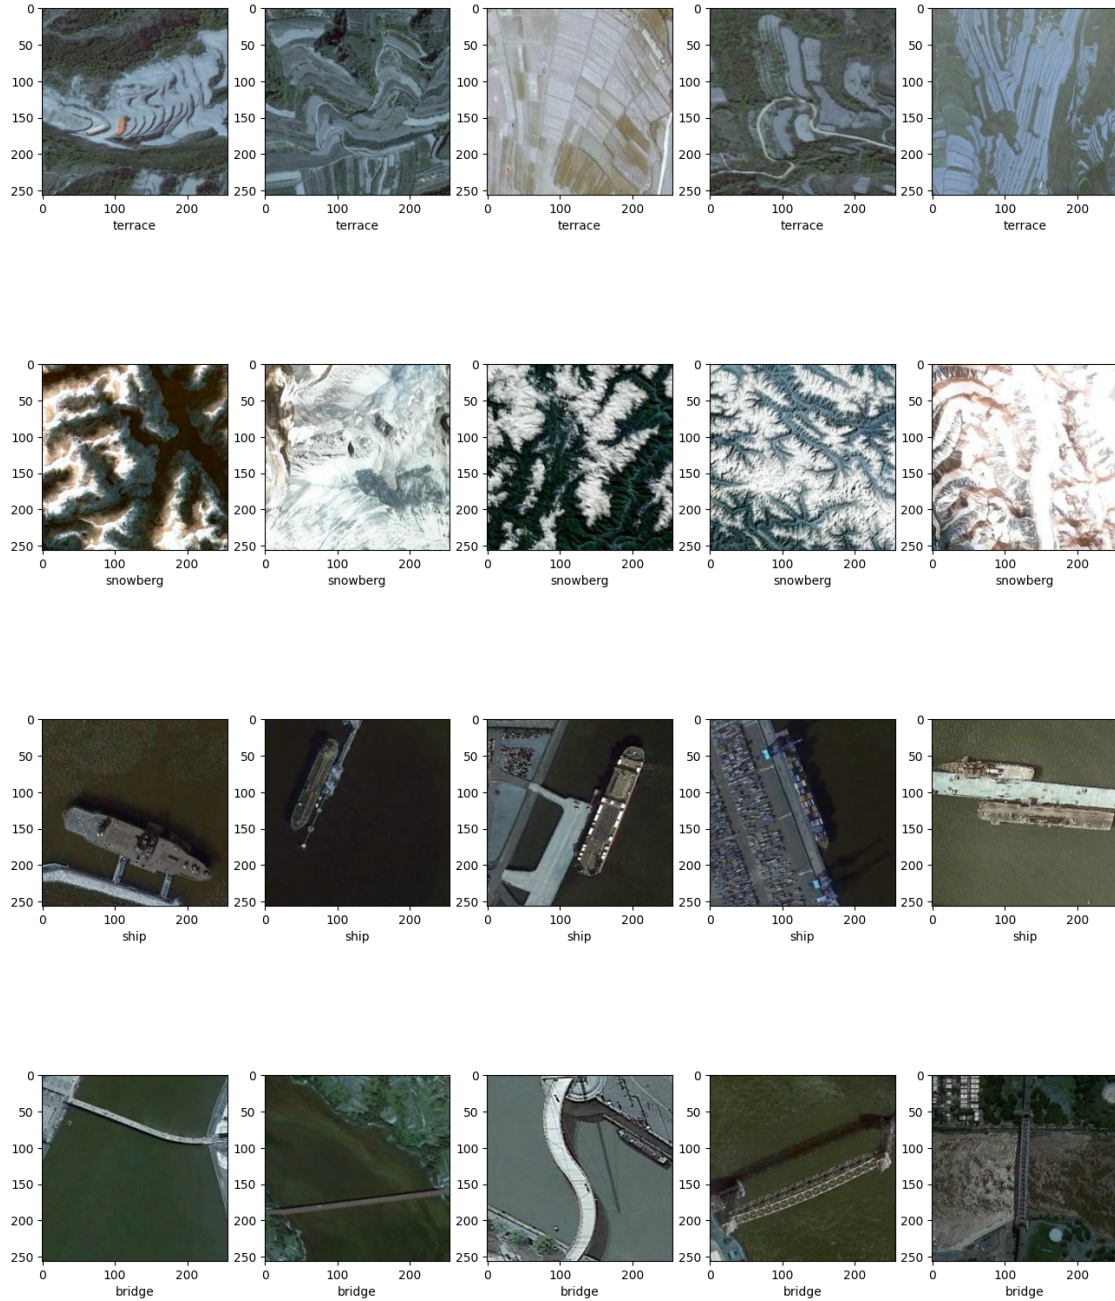
```

[ ]: random_labels = random.sample(LABELS, min(5, len(LABELS)))
for label in random_labels:
    airplane_dataset = data[data['label'] == label].imgpath
    img_airplane = airplane_dataset.tolist()
    p1 = []
    for img in img_airplane[:5]:
        x = cv2.imread(img)
        p1.append(x)

    plt.figure(figsize=(16,10))
    for i in range(1,6):
        plt.subplot(2, 5, i)
        plt.grid(False)
        plt.imshow(p1[i-1])
        plt.xlabel(label)
    plt.show()

```





Conclusion - Images have a resolution of 256 by 256 and 3 channels (RGB)

2.2 Partition dataset into folds

- add 5-fold for each sample (1,2,3,4,5): 1st 20%, 2nd 20%, etc.
- save current dataset on disk

```
[ ]: # - add label index for training
label_index = []
for lab in data.label:
    label_index.append(LABELS.index(lab))
data['label_index'] = label_index

[ ]: def partition_dataset(data, folds=5):
    # partition the dataset into folds
    partition_size = int(100 / folds)
    data = data.sample(frac=1).reset_index(drop=True) # resample dataset
    ↪randomly.
    folds = []
    n = len(data['label_index'])
    fold = 0
    for i in range(n):
        if ((i / n) * 100) % partition_size == 0: # folds are %20, for 80%
    ↪train and 20% val in 5-fold;
        fold += 1
        folds.append(fold)
    data['fold'] = folds
    data.sort_values(by=['label_index'], inplace=True) # sort values
    return data

[ ]: data = partition_dataset(data, folds=K_FOLDS)

[ ]: data.to_csv("dataset_resisc45.csv", index=False) # save dataset on disk
```

2.3 Data load

```
[ ]: df = pd.read_csv("dataset_resisc45.csv")

[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31500 entries, 0 to 31499
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  -
0   imgpath         31500 non-null  object
1   label           31500 non-null  object
2   label_index     31500 non-null  int64
3   fold            31500 non-null  int64
dtypes: int64(2), object(2)
memory usage: 984.5+ KB
```

```
[ ]: df["fold"].value_counts()
```

```
[ ]: fold
      2    6300
      3    6300
      1    6300
      5    6300
      4    6300
      Name: count, dtype: int64
```

2.4 Train dataset loader

```
[ ]: # Dataset loader
import enum

class DatasetTypes(enum.Enum):
    train = "train"
    val = "val"
    test = "test"

RES_X = 256
RES_Y = 256

class DatasetResisc45(Dataset):
    def __init__(self, dataset_file, dataset_type=None, val_fold=None,
        shuffle=False, transform=None, target_transform=None):
        df = pd.read_csv(dataset_file)
        if shuffle:
            df = df.sample(frac=1).reset_index(drop=True)
        folds = list(df["fold"].unique()) # get all folds
        if val_fold:
            if val_fold not in folds:
                raise Exception("Fold not found.")
            if dataset_type == DatasetTypes.train:
                df.drop(index=df[df["fold"] == val_fold].index, inplace=True)
            # drop the validation fold
            elif dataset_type == DatasetTypes.val:
                df = df[df["fold"] == val_fold] # keep only the validation
            # fold in dataset
        self.img_labels = df
        self.transform = transform
        self.target_transform = target_transform

    def __len__(self):
        return len(self.img_labels)

    def __getitem__(self, idx):
        img_path = self.img_labels.iloc[idx, 0]
```



```

        image = cv2.imread(img_path)
        image = cv2.resize(image, dsize=(RES_Y, RES_X))
        image = image - 128.5
        image = np.moveaxis(image, -1, 0).astype(np.float32) / 255.0 # move channels first (3 x RES_Y x RES_X)
        label_index = self.img_labels.iloc[idx, 2]
        label = np.zeros(len(LABELS), dtype=np.float32)
        label[label_index] = 1.0
        if self.transform:
            image = self.transform(image)
        if self.target_transform:
            label = self.target_transform(label)
        return image, label

```

3 Classification problem

- 5-fold cross validation
- one-vs-all
- softmax

```

[ ]: # find CUDA / MPS / CPU device
device = (
    "cuda"
    if torch.cuda.is_available()
    else "mps"
    if torch.backends.mps.is_available()
    else "cpu"
)
print(f"Using {device} device")

```

Using mps device

3.1 Neural Networks Algorithms

```

[ ]: from torch import nn

# Define model architecture
# need to implement dimensions checking!

# Define model
class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()
        self.linear_relu_stack = nn.Sequential(
            nn.Linear(RES_Y * RES_X * 3, 512),
            nn.ReLU(),

```

```

        nn.Linear(512, 512),
        nn.ReLU(),
        nn.Linear(512, len(LABELS))
    )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits

class SatAlexNet(nn.Module):
    def __init__(self, num_channels=3):
        super(SatAlexNet, self).__init__()

        # add 1x1 conv with stride 1
        # add 3x3 conv with stride 2 in first layer

        # 1. initialize first set of CONV => RELU => POOL layers
        self.conv1 = nn.Conv2d(in_channels=num_channels, out_channels=50,
        ↪kernel_size=(7, 7), stride=2)
        self.relu1 = nn.ReLU()
        self.maxpool1 = nn.MaxPool2d(kernel_size=(3, 3), stride=(2, 2))

        # 2. initialize second set of CONV => RELU => POOL layers
        self.conv2 = nn.Conv2d(in_channels=50, out_channels=100,
        ↪kernel_size=(3, 3))
        self.relu2 = nn.ReLU()
        self.maxpool2 = nn.MaxPool2d(kernel_size=(3, 3), stride=(2, 2))

        # 3. initialize second set of CONV => RELU => POOL layers
        self.conv3 = nn.Conv2d(in_channels=100, out_channels=200,
        ↪kernel_size=(3, 3))
        self.relu3 = nn.ReLU()
        self.maxpool3 = nn.MaxPool2d(kernel_size=(3, 3), stride=(2, 2))

        # 4. initialize second set of CONV => RELU => POOL layers
        self.conv4 = nn.Conv2d(in_channels=200, out_channels=200,
        ↪kernel_size=(3, 3))
        self.relu4 = nn.ReLU()
        self.maxpool4 = nn.MaxPool2d(kernel_size=(3, 3), stride=(2, 2))

        # 5. initialize second set of CONV => RELU => POOL layers
        self.conv5 = nn.Conv2d(in_channels=200, out_channels=100,
        ↪kernel_size=(3, 3))
        self.relu5 = nn.ReLU()
        self.maxpool5 = nn.MaxPool2d(kernel_size=(3, 3), stride=(2, 2))

```

```

self.flatten = nn.Flatten()

# initialize first (and only) set of FC => RELU layers
self.fc1 = nn.Linear(in_features=2500, out_features=len(LABELS))

def forward(self, x):
    # pass the input through our first set of CONV => RELU =>
    # POOL layers
    x = self.conv1(x)
    x = self.relu1(x)
    x = self.maxpool1(x)
    # pass the output from the previous layer through the second
    # set of CONV => RELU => POOL layers
    x = self.conv2(x)
    x = self.relu2(x)
    x = self.maxpool2(x)
    # pass the output from the previous layer through the second
    # set of CONV => RELU => POOL layers
    x = self.conv3(x)
    x = self.relu3(x)
    x = self.maxpool3(x)
    # pass the output from the previous layer through the second
    # set of CONV => RELU => POOL layers
    x = self.conv4(x)
    x = self.relu4(x)
    x = self.maxpool4(x)
    # pass the output from the previous layer through the second
    # set of CONV => RELU => POOL layers
    x = self.conv5(x)
    x = self.relu5(x)
    x = self.maxpool5(x)
    # flatten the output from the previous layer and pass it
    # through our only set of FC => RELU layers
    x = self.flatten(x)
    x = self.fc1(x)
    return x

```

```

[ ]: # create model
def create_model_on_device(model_class, *args, **kwargs):
    print(f"device: {device}")
    model = model_class(*args, **kwargs).to(device) # create model on device
    print(str(model))
    return model

```

```

[ ]: demo_model = create_model_on_device(SatAlexNet, num_channels=3)

```

```

device: mps
SatAlexNet(
  (conv1): Conv2d(3, 50, kernel_size=(3, 3), stride=(1, 1))
  (relu1): ReLU()
  (maxpool1): MaxPool2d(kernel_size=(3, 3), stride=(2, 2), padding=0,
dilation=1, ceil_mode=False)
  (conv2): Conv2d(50, 100, kernel_size=(3, 3), stride=(1, 1))
  (relu2): ReLU()
  (maxpool2): MaxPool2d(kernel_size=(3, 3), stride=(2, 2), padding=0,
dilation=1, ceil_mode=False)
  (conv3): Conv2d(100, 200, kernel_size=(3, 3), stride=(1, 1))
  (relu3): ReLU()
  (maxpool3): MaxPool2d(kernel_size=(3, 3), stride=(2, 2), padding=0,
dilation=1, ceil_mode=False)
  (conv4): Conv2d(200, 200, kernel_size=(3, 3), stride=(1, 1))
  (relu4): ReLU()
  (maxpool4): MaxPool2d(kernel_size=(3, 3), stride=(2, 2), padding=0,
dilation=1, ceil_mode=False)
  (conv5): Conv2d(200, 100, kernel_size=(3, 3), stride=(1, 1))
  (relu5): ReLU()
  (maxpool5): MaxPool2d(kernel_size=(3, 3), stride=(2, 2), padding=0,
dilation=1, ceil_mode=False)
  (flatten): Flatten(start_dim=1, end_dim=-1)
  (fc1): Linear(in_features=2500, out_features=45, bias=True)
)

```

```

[ ]: loss_fn = nn.CrossEntropyLoss()

def train_one_epoch(dataloader, model):
    optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)

    print("Started train.")
    size = len(dataloader.dataset)

    model.train() # set model to train mode

    for batch, (X, y) in enumerate(dataloader):
        X, y = X.to(device), y.to(device)

        # Compute prediction error
        pred = model(X)
        loss = loss_fn(pred, y)

        # Backpropagation
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

```

```

        if batch % 100 == 0:
            loss, current = loss.item(), (batch + 1) * len(X)
            print(f"loss: {loss:>7f} [{current:>5d}/{size:>5d}]")

def val_one_epoch(dataloader, model):
    print("Started validation.")
    size = len(dataloader.dataset)
    num_batches = len(dataloader)

    model.eval() # set model to evaluation mode

    test_loss, correct = 0, 0
    with torch.no_grad():
        for X, y in dataloader:
            X, y = X.to(device), y.to(device)
            pred = model(X)
            test_loss += loss_fn(pred, y).item()
            error = (nn.Softmax(dim=1)(pred).argmax(1) == y.argmax(1)).
↪type(torch.float)
            correct += error.sum().item()
    test_loss /= num_batches
    correct /= size
    print(f"Test Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss:↪
↪{test_loss:>8f} \n")

```

3.2 5 Fold train validation

- trebuie sa reducem timpii de antrenament

```

[ ]: def kfold_train():
    epochs = 5
    batch_size = 64

    for val_fold in range(K_FOLDS):
        print(f"Fold: {val_fold}\n-----")
        model = create_model_on_device(SatAlexNet)
        train_set = DatasetResisc45("dataset_resisc45.csv",↪
↪dataset_type=DatasetTypes.train, val_fold=val_fold)
        val_set = DatasetResisc45("dataset_resisc45.csv",↪
↪dataset_type=DatasetTypes.val, val_fold=val_fold)
        train_dataloader = DataLoader(train_set, batch_size=batch_size,↪
↪shuffle=True)
        val_dataloader = DataLoader(val_set, batch_size=batch_size,↪
↪shuffle=False)

        for t in range(epochs):

```

```

print(f"Epoch {t+1}\n-----")
train_one_epoch(train_dataloader, model)
val_one_epoch(val_dataloader, model)

```

```

[ ]: def simple_train():
    model = create_model_on_device(SatAlexNet)
    train_set = DatasetResisc45("dataset_resisc45.csv",
    ↪dataset_type=DatasetTypes.train, val_fold=5)
    val_set = DatasetResisc45("dataset_resisc45.csv", dataset_type=DatasetTypes.
    ↪val, val_fold=5)
    train_dataloader = DataLoader(train_set, batch_size=64, shuffle=True)
    val_dataloader = DataLoader(val_set, batch_size=64, shuffle=False)

    epochs = 10
    for t in range(epochs):
        print(f"Epoch {t+1}\n-----")
        train_one_epoch(train_dataloader, model)
        val_one_epoch(val_dataloader, model)

```

```

[ ]: simple_train()

```

```

device: mps
SatAlexNet(
  (conv1): Conv2d(3, 50, kernel_size=(3, 3), stride=(1, 1))
  (relu1): ReLU()
  (maxpool1): MaxPool2d(kernel_size=(3, 3), stride=(2, 2), padding=0,
dilation=1, ceil_mode=False)
  (conv2): Conv2d(50, 100, kernel_size=(3, 3), stride=(1, 1))
  (relu2): ReLU()
  (maxpool2): MaxPool2d(kernel_size=(3, 3), stride=(2, 2), padding=0,
dilation=1, ceil_mode=False)
  (conv3): Conv2d(100, 200, kernel_size=(3, 3), stride=(1, 1))
  (relu3): ReLU()
  (maxpool3): MaxPool2d(kernel_size=(3, 3), stride=(2, 2), padding=0,
dilation=1, ceil_mode=False)
  (conv4): Conv2d(200, 200, kernel_size=(3, 3), stride=(1, 1))
  (relu4): ReLU()
  (maxpool4): MaxPool2d(kernel_size=(3, 3), stride=(2, 2), padding=0,
dilation=1, ceil_mode=False)
  (conv5): Conv2d(200, 100, kernel_size=(3, 3), stride=(1, 1))
  (relu5): ReLU()
  (maxpool5): MaxPool2d(kernel_size=(3, 3), stride=(2, 2), padding=0,
dilation=1, ceil_mode=False)
  (flatten): Flatten(start_dim=1, end_dim=-1)
  (fc1): Linear(in_features=2500, out_features=45, bias=True)
)
Epoch 1

```

```
-----  
Started train.  
loss: 3.806273 [ 64/25200]  
loss: 3.474065 [ 6464/25200]  
loss: 2.936102 [12864/25200]  
loss: 2.446790 [19264/25200]  
Started validation.  
Test Error:  
Accuracy: 32.9%, Avg loss: 2.364211
```

Epoch 2

```
-----  
Started train.  
loss: 2.398020 [ 64/25200]  
loss: 1.911188 [ 6464/25200]  
loss: 2.067113 [12864/25200]  
loss: 2.186711 [19264/25200]  
Started validation.  
Test Error:  
Accuracy: 43.1%, Avg loss: 1.932437
```

Epoch 3

```
-----  
Started train.  
loss: 1.844561 [ 64/25200]  
loss: 2.025104 [ 6464/25200]  
loss: 1.833625 [12864/25200]  
loss: 1.731433 [19264/25200]  
Started validation.  
Test Error:  
Accuracy: 48.9%, Avg loss: 1.700356
```

Epoch 4

```
-----  
Started train.  
loss: 1.305466 [ 64/25200]  
loss: 1.348948 [ 6464/25200]  
loss: 1.635152 [12864/25200]  
loss: 1.632333 [19264/25200]  
Started validation.  
Test Error:  
Accuracy: 52.2%, Avg loss: 1.583759
```

Epoch 5

```
-----  
Started train.  
loss: 1.613610 [ 64/25200]  
loss: 1.260794 [ 6464/25200]
```

loss: 1.493123 [12864/25200]
loss: 1.378588 [19264/25200]
Started validation.
Test Error:
Accuracy: 52.8%, Avg loss: 1.607404

Epoch 6

Started train.
loss: 1.569013 [64/25200]
loss: 1.201636 [6464/25200]
loss: 1.124559 [12864/25200]
loss: 1.296306 [19264/25200]
Started validation.
Test Error:
Accuracy: 57.2%, Avg loss: 1.433341

Epoch 7

Started train.
loss: 1.261605 [64/25200]
loss: 1.517849 [6464/25200]
loss: 1.233839 [12864/25200]
loss: 1.598638 [19264/25200]
Started validation.
Test Error:
Accuracy: 61.7%, Avg loss: 1.299462

Epoch 8

Started train.
loss: 1.280254 [64/25200]
loss: 0.906131 [6464/25200]
loss: 1.044051 [12864/25200]
loss: 0.985112 [19264/25200]
Started validation.
Test Error:
Accuracy: 61.4%, Avg loss: 1.303278

Epoch 9

Started train.
loss: 0.748255 [64/25200]
loss: 1.170680 [6464/25200]
loss: 1.078530 [12864/25200]
loss: 0.927029 [19264/25200]
Started validation.
Test Error:

Accuracy: 63.5%, Avg loss: 1.240218

Epoch 10

Started train.

loss: 1.012814 [64/25200]

loss: 1.044277 [6464/25200]

loss: 1.018787 [12864/25200]

loss: 1.054679 [19264/25200]

Started validation.

Test Error:

Accuracy: 63.7%, Avg loss: 1.231658

3.3 Bias-variance trade

- work in progress..

```
[ ]: %matplotlib ipynb

epochs = np.arange(0, 20, 1)
train_losses = []
val_losses = []
for epoc in epochs:
    train_losses.append(random.random())
    val_losses.append(random.random())

x = epochs
y = train_losses
y2 = val_losses

fig, ax = plt.subplots()
ax.plot(x, y)
ax.plot(x, y2)
#specify axis tick step sizes
_ = plt.xticks(np.arange(min(x), max(x)+1, 1))
_ = plt.yticks(np.arange(0, max(y)+0.1, 0.1))

ax.set_title("train loss")
plt.show()
```



[]: